

Factor Analysis

Max Turgeon

STAT 4690—Applied Multivariate Analysis

Latent variable models

- With PCA, we saw how we could reduce the dimension of data using the eigenvectors of the sample covariance matrix.
- Conversely, we could construe PCA has a generative model, where the principal components give rise to the observed data.
- **Latent Variable Models** formalise this idea:
 - Latent (i.e. unobserved) variables \mathbf{F} give rise to observed data \mathbf{Y} through a *specified* model.

Factor Analysis i

- **Factor Analysis** is a special kind of latent variable model.
- Let \mathbf{Y} be a p -dimensional vector with mean μ and covariance matrix Σ .
- Let \mathbf{F} be a m -dimensional *latent* vector.
- The *orthogonal factor model* is given by

$$\mathbf{Y} - \mu = L\mathbf{F} + \mathbf{E},$$

where L is a $p \times m$ *matrix of factor loadings*, and \mathbf{E} is a p -dimensional vector of *errors*.

Factor Analysis ii

- F are also called *common factors*; E are also called *specific factors*.
- **Note:** This is essentially a multivariate regression model, but where the covariates are unobserved.

Assumptions i

- The model above is generally not identifiable, since there are too many parameters.
- We therefore need to impose further restrictions:
 - $E(\mathbf{F}) = 0$
 - $\text{Cov}(\mathbf{F}) = I$
 - $E(\mathbf{E}) = 0$
 - $\text{Cov}(\mathbf{E}) = \Psi = \begin{pmatrix} \psi_1 & 0 & \cdots & 0 \\ 0 & \psi_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \psi_p \end{pmatrix}$
 - $\text{Cov}(\mathbf{F}, \mathbf{E}) = 0$

Assumptions ii

- In other words:
- Both common and specific factors have mean zero;
- They are uncorrelated;
- The common factors are mutually uncorrelated and standardised;
- The specific factors each affect only one observed variable.

Structured Covariance i

- As a consequence of these assumptions, we can derive an assumption on the structure of $\Sigma = \text{Cov}(\mathbf{Y})$:

$$\begin{aligned}\Sigma &= E\left((\mathbf{Y} - \mu)(\mathbf{Y} - \mu)^T\right) \\ &= E\left((L\mathbf{F} + \mathbf{E})(L\mathbf{F} + \mathbf{E})^T\right) \\ &= LE(\mathbf{F}\mathbf{F}^T)L + E(\mathbf{E}\mathbf{F}^T)L^T + LE(\mathbf{F}\mathbf{E}^T) + E(\mathbf{E}\mathbf{E}^T) \\ &= LIL^T + 0L^T + L0 + \Psi \\ &= LL^T + \Psi.\end{aligned}$$

Structured Covariance ii

- Similarly, we can show that

$$\text{Cov}(\mathbf{Y}, \mathbf{F}) = L.$$

- If we write ℓ_{ij} for the (i, j) -th element of L , we see that

$$\text{Var}(Y_i) = \sum_{k=1}^m \ell_{ik}^2 + \psi_i.$$

- Crucially, these assumptions are **testable**. In other words, we can check whether they are reasonable for our data.

Example i

- Let's look at an example where there is no solution.
- Assume $p = 3$, $m = 1$, with

$$\Sigma = \begin{pmatrix} 1 & 0.9 & 0.7 \\ 0.9 & 1 & 0.4 \\ 0.7 & 0.4 & 1 \end{pmatrix}.$$

- From our assumptions on the covariance structure, we derive a system of equations

$$\begin{aligned} 1 &= \ell_{11}^2 + \psi_1 & 0.9 &= \ell_{11}\ell_{21} & 0.7 &= \ell_{11}\ell_{31} \\ & & 1 &= \ell_{22}^2 + \psi_2 & 0.4 &= \ell_{21}\ell_{31} \\ & & & & 1 &= \ell_{33}^2 + \psi_3 \end{aligned}$$

Example ii

- From $0.7 = \ell_{11}\ell_{31}$ and $0.4 = \ell_{21}\ell_{31}$, we get

$$\ell_{21} = \frac{0.4}{0.7}\ell_{11}.$$

- But since $0.9 = \ell_{11}\ell_{21}$, we can conclude that

$$\ell_{11} = \pm 1.255.$$

- However, since the first component Y_1 has unit variance, $\ell_{11} = \text{Corr}(Y_1, F_1)$, and therefore the correlation is out of bounds.
- Similarly, we get

$$\psi_1 = 1 - \ell_{11}^2 = 1 - 1.575 = -0.575.$$

Example iii

- But since ψ_1 is the variance of the first error term, we once again get a non-sensical solution.

Factor Rotation i

- Even with our assumptions above, our model is still not uniquely identified.
- Let T be an $m \times m$ orthogonal matrix. We have

$$\begin{aligned}\mathbf{Y} - \mu &= L\mathbf{F} + \mathbf{E} \\ &= LTT^T\mathbf{F} + \mathbf{E} \\ &= \tilde{L}\tilde{\mathbf{F}} + \mathbf{E},\end{aligned}$$

where $\tilde{L} = LT$ and $\tilde{\mathbf{F}} = T^T\mathbf{F}$.

Factor Rotation ii

- Both models lead to the same covariance matrix:

$$\Sigma = LL^T + \Psi = LTT^TL^T + \Psi = \tilde{L}\tilde{L}^T + \Psi.$$

- As we will see, this will turn out to be a blessing in disguise:
 - We will impose a uniqueness condition to get one solution.
 - Then we will rotate our solution using T to improve interpretation.

Estimation—Principal Component Method i

- Recall the spectral decomposition of the covariance matrix:

$$\Sigma = \sum_{i=1}^p \lambda_i w_i w_i^T,$$

with $\lambda_1 \geq \dots \geq \lambda_p$.

- If we let W be the matrix whose i -th column is $\sqrt{\lambda_i} w_i$, we can rewrite the spectral decomposition as

$$\Sigma = WW^T.$$

- In other words, if we let $m = p$ and $\Psi = 0$, we see that we recover the orthogonal factor model with $L = W$.

Estimation—Principal Component Method ii

- Of course, this is not very satisfactory, as the dimension of the common factors is the same as that of the original data.
- Instead, we select $m < p$ using one of the methods we discussed with PCA and we approximate

$$\Sigma \approx \sum_{i=1}^m \lambda_i w_i w_i^T.$$

- If we let L be the $p \times m$ matrix whose i -th column is $\sqrt{\lambda_i} w_i$, we can estimate Ψ as follows:

$$\psi_i = \sigma_{ii} - \sum_{j=1}^m \ell_{ij}^2.$$

Algorithm

1. Let $\hat{\lambda}_1 \cdots > \hat{\lambda}_p$ and $\hat{w}_1, \dots, \hat{w}_p$ be the eigenvalues and eigenvectors of the covariance matrix S_n .
2. Select m using one of the PCA criteria.
3. Estimate \hat{L} with the $p \times m$ matrix whose i -th column is $\sqrt{\hat{\lambda}_i} \hat{w}_i$.
4. Estimate $\hat{\Psi}$ with the diagonal elements of $S_n - \hat{L}\hat{L}^T$.

Example i

- We are going to use the `bfi` dataset in the `psych` package.
- Contains data on 25 personality items grouped in 5 categories:
 - A (Agreeableness)
 - C (Conscientiousness)
 - E (Extraversion)
 - N (Neuroticism)
 - O (Openness)

Example ii

```
library(psych)
```

```
dim(bfi)
```

```
## [1] 2800    28
```

```
tail(names(bfi), n = 3)
```

```
## [1] "gender"      "education" "age"
```

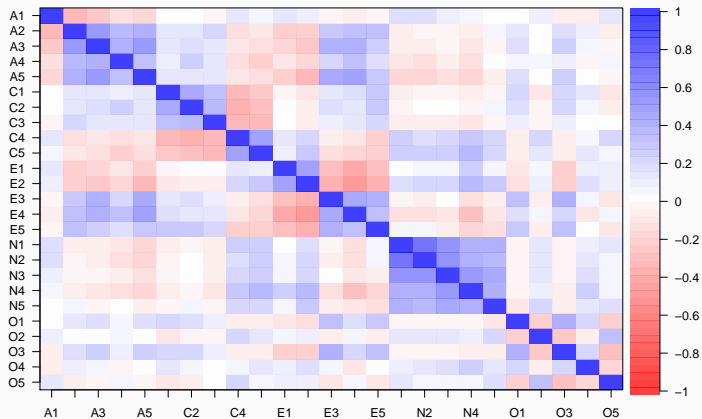
Example iii

```
library(tidyverse)

# Remove demographic variable and keep complete data
data <- bfi %>%
  select(-gender, -education, -age) %>%
  filter(complete.cases(.))

cor.plot(cor(data))
```

Example iv



Example v

```
decomp <- prcomp(data)
summary(decomp)$importance[,1:3]
```

##	PC1	PC2	PC3
## Standard deviation	3.291635	2.451538	2.030393
## Proportion of Variance	0.215650	0.119620	0.082050
## Cumulative Proportion	0.215650	0.335270	0.417320

Example vi

```
cum_prop <- decomp %>%  
  summary %>%  
  .[["importance"]] %>%  
  .["Cumulative Proportion",]
```

```
which(cum_prop > 0.8)[1]
```

```
## PC14
```

```
## 14
```

Example vii

```
Lhat <- decomp$rotation[,1:14] %*%  
  diag(decomp$sdev[1:14])  
Psi_hat <- diag(cov(data) - tcrossprod(Lhat))  
  
# Sum squared error  
sum((cov(data) - tcrossprod(Lhat) - diag(Psi_hat))^2)  
  
## [1] 3.645694  
  
# Compare to the total variance  
sum(diag(cov(data)))
```

Example viii

```
## [1] 50.24287
```

```
# Our FA model explains:
```

```
sum(colSums(Lhat^2)/sum(diag(cov(data))))
```

```
## [1] 0.8160488
```


Comments

- In our example above, we saw that 14 factors explained 82% of the total variance.
- Two common default values for m in statistical softwares:
 - Number of positive eigenvalues of the sample covariance matrix.
 - Number of eigenvalues greater than one for the sample correlation matrix.
- In our example, the first criterion would lead to $m = p$, which is not very helpful

Example (cont'd) i

```
(m <- sum(eigen(cor(data))$values > 1))
```

```
## [1] 6
```

```
Lhat <- decomp$rotation[,seq_len(m)] %*%  
  diag(decomp$sdev[seq_len(m)])
```

```
Psi_hat <- diag(cov(data) - tcrossprod(Lhat))
```

```
# Sum squared error
```

```
sum((cov(data) - tcrossprod(Lhat) - diag(Psi_hat))^2)
```

Example (cont'd) ii

```
## [1] 6.999035
```

```
# Compare to the total variance  
sum(diag(cov(data)))
```

```
## [1] 50.24287
```

```
# Our FA model explains:  
sum(colSums(Lhat2)/sum(diag(cov(data))))
```

```
## [1] 0.5910586
```

Example (cont'd) iii

We can also visualize the fit


```
Sn <- cov(data)
```

```
Sn_fit <- tcrossprod(Lhat) - diag(Psi_hat)
```

```
library(lattice)
```

```
levelplot(Sn - Sn_fit)
```

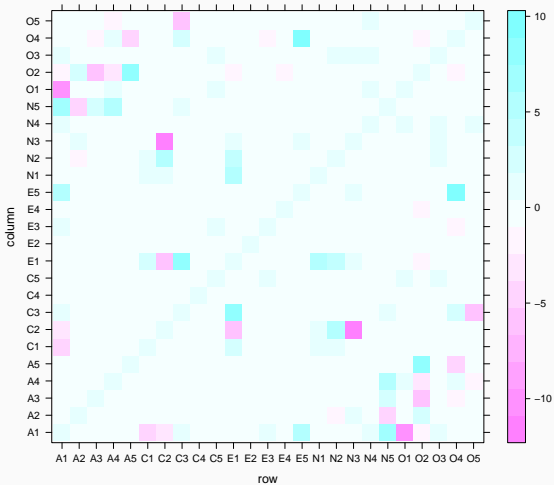
Example (cont'd) iv



Example (cont'd) v

```
# Or if you prefer % difference  
levelplot((Sn - Sn_fit)/Sn)
```

Example (cont'd) vi



Estimation—Maximum Likelihood Method i

- This estimation method assumes that both common factors \mathbf{F} and specific factors \mathbf{E} follow a multivariate normal distribution
 - $\mathbf{F} \sim N_m(0, I)$
 - $\mathbf{E} \sim N_p(0, \Psi)$
- From this assumption, it follows that \mathbf{Y} also follows a multivariate normal distribution
 - $\mathbf{Y} \sim N_p(\mu, LL^T + \Psi)$
- Therefore, we can write down the likelihood in terms of both L and Ψ .

Estimation—Maximum Likelihood Method ii

- However, because of the factor rotation problem, we need to impose a constraint in order to obtain a unique solution:
 - $L^T \Psi^{-1} L = \Delta$ is diagonal.
- Given this assumption, the maximum likelihood estimates \hat{L} and $\hat{\Psi}$ can be found using an iterative algorithm.
- We will not go into the details of the algorithm (but see Johnson & Wichern, Supplement 9A if interested).
 - Instead, we will rely on the R function `stats::factanal`.

Example (cont'd) i

CAREFUL: uses correlation matrix

```
fa_decomp <- factanal(data, factors = m,  
                      rotation = 'none')
```

*# Uniquenesses are the diagonal elements
of the matrix Psi*

```
Psi_mle <- fa_decomp$uniquenesses  
Lmle <- fa_decomp$loadings
```

Example (cont'd) ii

We get an estimate of the correlation

```
R_mle <- tcrossprod(Lmle) - diag(Psi_mle)
```

```
sum((cor(data) - R_mle)^2)
```

```
## [1] 31.97525
```

Our FA model explains:

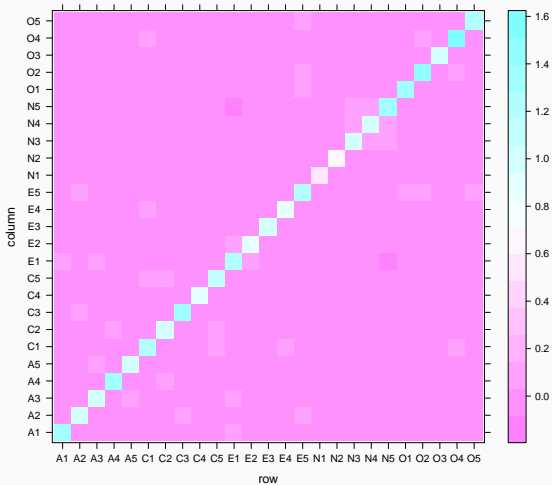
```
sum(colSums(Lmle^2)/ncol(data))
```

```
## [1] 0.4498342
```

Example (cont'd) iii

```
levelplot(cor(data) ~ R_mle)
```

Example (cont'd) iv



Example (cont'd) v

```
# To factor the covariance matrix
```

```
# Use psych::fa
```

```
fa_decomp <- psych::fa(data, nfactors = m,  
                        rotate = "none",  
                        covar = TRUE,  
                        fm = "ml")
```

```
# Extract estimates
```

```
Psi_mle <- fa_decomp$uniquenesses
```

```
Lmle <- fa_decomp$loadings
```

Example (cont'd) vi

We get an estimate of the covariance

```
Sn_mle <- tcrossprod(Lmle) - diag(Psi_mle)
```

```
sum((Sn - Sn_mle)^2)
```

```
## [1] 128.025
```

Our FA model explains:

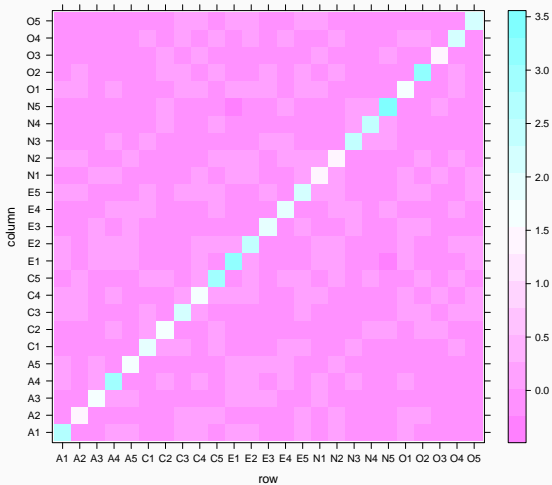
```
sum(colSums(Lmle^2)/sum(diag(Sn)))
```

```
## [1] 0.459665
```

Example (cont'd) vii

```
levelplot(Sn ~ Sn_mle)
```



Example (cont'd) viii



Example (cont'd) ix

```
# Compare MLE with PC estimate  
levelplot((Sn_fit - Sn_mle)/Sn)
```

Example (cont'd) x



Comments about estimation i

- There are other methods of estimating the loadings and the “uniquenesses”
 - Ordinary Least Squares
 - Weighted OLS
 - Principal factor
- With so many choices of estimation methods, it can be hard to compare statistical softwares
 - And you also have to read the manual in order to know what is going on...
- You also always have the choice between factoring the covariance or the correlation matrix

Comments about estimation ii

- Which one you choose depends on the commensurability of your variables (just like in PCA)
- Finally, it's always a good idea to compare the output of multiple estimation strategies
 - If your model is a good fit, you should get a similar answer regardless of the method.

Factor Rotation Redux i

- As we saw earlier, any orthogonal matrix T gives rise to the same factor analysis model

$$\Sigma = LL^T + \Psi = LTT^TL^T + \Psi = \tilde{L}\tilde{L}^T + \Psi.$$

- In other words, we cannot choose T to maximise the goodness of fit.
 - We need another criterion
- Intuitively, to ease interpretation, we want each variable to have large loadings for one factor and negligible loadings for the other ones.

Factor Rotation Redux ii

- https://maxturgeon.ca/f19-stat4690/factor_rotation.gif
- One common analytic criterion that formalises this idea is the **varimax criterion**.
 - Resulting loadings are called *varimax loadings*
- We have

$$\text{VARIMAX} \propto \sum_{j=1}^m \left(\begin{array}{c} \text{Variance of squares of scales loadings} \\ \text{for } j\text{-th factor} \end{array} \right)$$

- More precisely:
 - Let $\tilde{\ell}_{ij}$ be the (i, j) -th entry of the matrix $\tilde{L} = LT$. In other words, $\tilde{\ell}_{ij}$ depends on T .

Factor Rotation Redux iii

- Let $\tilde{h}_i^2 = \sum_{j=1}^m \tilde{\ell}_{ij}^2$.
- Define the scaled loadings $\tilde{\ell}_{ij}^* = \tilde{\ell}_{ij}/h_i$.
- The varimax criterion V is given as



$$V = \frac{1}{p} \sum_{j=1}^m \left(\sum_{i=1}^p \tilde{\ell}_{ij}^{*4} - \frac{1}{p} \left(\sum_{i=1}^p \tilde{\ell}_{ij}^{*2} \right)^2 \right).$$

- In R, you can compute the rotated loadings using the `stats::varimax` function. Alternatively, the function `stats::factanal` can compute the rotation for you as part of the factor analysis (and so does `psych::fa`).

Example (cont'd) i

```
# Let's start with m=2 for visualization  
fa_decomp <- factanal(data, factors = 2,  
                      rotation = 'none')  
  
initial_loadings <- fa_decomp$loadings  
varimax_loadings <- varimax(initial_loadings)
```

Example (cont'd) ii



Example (cont'd) iii

You can extract the matrix T

```
varimax_loadings$rotmat
```

```
##           [,1]      [,2]
```

```
## [1,] 0.7810310 -0.6244923
```

```
## [2,] 0.6244923  0.7810310
```

We can also get the angle of rotation

```
acos(varimax_loadings$rotmat[1,1])
```

```
## [1] 0.6744813
```

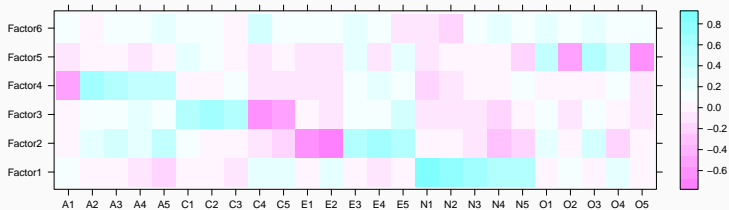
Example (cont'd) iv

```
# In more dimensions
```

```
fa_decomp <- factanal(data, factors = m,  
                      rotation = 'varimax')
```

```
levelplot(unclass(fa_decomp$loadings),  
          xlab = "", ylab = "")
```

Example (cont'd) v



Comments

- As with estimation, there are many more rotation methods.
 - See for example the help page
`?GPArotation::rotations`
- One particular class of rotations are called *oblique*
 - The matrix T is no longer constrained to be orthogonal.
- Factor rotation is especially useful with loadings obtained through MLE
 - Recall the constraint on $L^T \Psi^{-1} L$ being diagonal
- Factor rotations are also sometimes used with PCA loadings.