

Manifold Learning

Max Turgeon

STAT 4690—Applied Multivariate Analysis

Dimension reduction redux i

- Recall Pearson's approach to PCA: **best approximation of the data by a linear manifold.**
- Let's unpack this definition:
 - We're looking for a linear subspace of \mathbb{R}^p of dimension k .
 - For a fixed k , we want to minimise the error when projecting onto the linear subspace.
 - We can also identify that subspace with \mathbb{R}^k (e.g. for visualisation).

Dimension reduction redux ii

- **Manifold learning** is a nonlinear approach to dimension reduction, where:
 - We assume the data lies on (or close to) a nonlinear manifold of dimension k in \mathbb{R}^p .
 - We project the data from the manifold to \mathbb{R}^k .
- There are two main classes of methods:
 - Distance preserving (e.g. Isomap);
 - Topology preserving (e.g. Locally linear embedding)

Manifolds–Definition

- Roughly speaking, **manifolds** of dimension k are geometric objects that locally look like \mathbb{R}^k .
 - Every point on the manifold has an open neighbourhood that is equivalent to an open ball in \mathbb{R}^k .
- Examples in \mathbb{R}^p include any curve, the $(p - 1)$ -dimensional sphere, or any linear subspace.
- Some manifolds have boundaries (e.g. a cylinder) or corners (e.g. a cube).

Swiss roll i



```
n <- 1000  
F1 <- runif(n, 0, 10)  
F2 <- runif(n, -1, 1)
```

```
X <- F1 * cos(F1)  
Y <- F2  
Z <- F1 * sin(F1)
```

Swiss roll ii

```
library(scatterplot3d)
library(colorspace)
colours <- cut(F1, breaks = seq(0, 10),
               labels = diverging_hcl(10))
par(mfrow = c(1, 2))
scatterplot3d(X, Y, Z, pch = 19, asp = 1,
               color = colours)
scatterplot3d(X, Y, Z, pch = 19, asp = 1,
               color = colours, angle = 80)
```

Swiss roll iii




Let's see if PCA can unroll the Swiss roll

```
decomp <- prcomp(cbind(X, Y, Z))
```

```
plot(decomp$x[,1:2],  
     col = as.character(colours), pch = 19)
```


Swiss roll v



MNIST data revisited

- To study the nonlinear dimension reduction methods in this lecture, we will restrict our attention to the digit 2 in the MNIST dataset.
- The reason: we can think of the different shapes of 2 as “smooth deformations” of one another.
 - This would work for other digits too (e.g. 6, 9, 8), but not all (e.g. 4, 7).

Example i

```
library(dslabs)
library(tidyverse)


mnist <- read_mnist()

data <- mnist$train$images[mnist$train$labels == 2, ]
```

Example ii

```
par(mfrow = c(1, 2))  
# With crossing  
matrix(data[1,], ncol = 28)[ , 28:1] %>%  
  image(col = gray.colors(12, rev = TRUE),  
        axes = FALSE, asp = 1)  
# Without crossing  
matrix(data[4,], ncol = 28)[ , 28:1] %>%  
  image(col = gray.colors(12, rev = TRUE),  
        axes = FALSE, asp = 1)
```


Example iii



Example iv

```
decomp <- prcomp(data)
decomp$x[,1:2] %>%
  as.data.frame() %>%
  ggplot(aes(PC1, PC2)) +
  geom_point(alpha = 0.5) +
  theme_minimal()
```



Example v



Example vi

```
# First PC
par(mfrow = c(1, 2))
index_right <- which.max(decomp$x[,1])
matrix(data[index_right,], ncol = 28)[ , 28:1] %>%
  image(col = gray.colors(12, rev = TRUE),
        axes = FALSE, asp = 1)
index_left <- which.min(decomp$x[,1])
matrix(data[index_left,], ncol = 28)[ , 28:1] %>%
  image(col = gray.colors(12, rev = TRUE),
        axes = FALSE, asp = 1)
```



Example vii



Example viii

```
# Second PC
par(mfrow = c(1, 2))
index_top <- which.max(decomp$x[,2])
matrix(data[index_top,], ncol = 28)[ , 28:1] %>%
  image(col = gray.colors(12, rev = TRUE),
        axes = FALSE, asp = 1)
index_bottom <- which.min(decomp$x[,2])
matrix(data[index_bottom,], ncol = 28)[ , 28:1] %>%
  image(col = gray.colors(12, rev = TRUE),
        axes = FALSE, asp = 1)
```

Example ix



Example x

PC1=-1446



PC1=-906



PC1=-718



PC1=-546



PC1=-415



PC1=-290



PC1=-170



PC1=-57



PC1=89



PC1=242



PC1=409



PC1=597



PC1=773



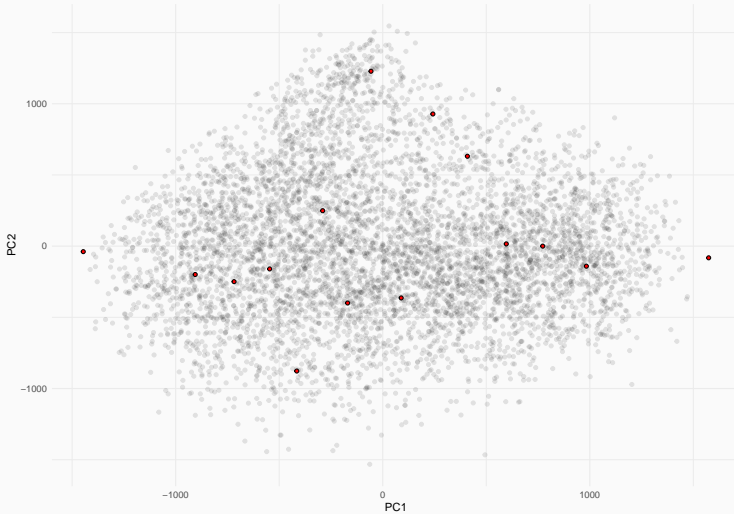
PC1=983



PC1=1574



Example xi



Isomap

- Let's look at the algorithm and study each step separately.

Basic algorithm

1. Create a graph \mathcal{G} from the data, where each data point is a node, and two nodes are connected if they are “neighbours”.
2. Each edge gets a weight corresponding to the Euclidean distance between the two data points.
3. Create a distance matrix Δ , where the (i, j) -th element is the length of the shortest path in \mathcal{G} between the data points corresponding to nodes i and j .
4. Perform metric Multidimensional Scaling on Δ to obtain the projection onto a lower dimensional subspace.

Definition of neighbourhood

- Two ways of defining the neighbours of a point \mathbf{Y} :
 - For an integer $K \geq 1$, we could look at the K -nearest neighbours, i.e. the K points $\mathbf{Y}_1, \dots, \mathbf{Y}_K$ that are closest (in Euclidean distance) to \mathbf{Y} .
 - For a real number $\epsilon > 0$, we could look at all points $\mathbf{Y}_1, \dots, \mathbf{Y}_{n(\epsilon)}$ whose distance from \mathbf{Y} is less than ϵ .
- Note:** The first definition guarantees that every point has neighbours, whereas you could get unconnected points using the second definition.
- You could also use a hybrid of both approaches where you take the K -nearest neighbours, but discard neighbours that are “too far away”.

Shortest path distance i

- Once we have our weighted graph \mathcal{G} (i.e. nodes represent data points, edges represent neighbours, weights are Euclidean distances), we can compute the length of any path from \mathbf{Y}_i to \mathbf{Y}_j by summing the weights of all the edges along the path.
- We then define a **distance** function on \mathcal{G} by

$$\Delta_{ij} = \min \{ \text{Length of path } \gamma \mid \gamma \text{ is a path from } \mathbf{Y}_i \text{ to } \mathbf{Y}_j \}.$$

Shortest path distance ii

- There are efficient algorithms for computing this distance for any weighted graph:
 - Dijkstra's algorithm;
 - Floyd–Warshall algorithm.
- For more details about these algorithms, take a course on graph theory!

Multidimensional Scaling

Recall the algorithm for MDS.

Algorithm (MDS)

Input: Δ ; Output: \tilde{X}

1. Create the matrix D containing the square of the entries in Δ .
2. Create S by centering both the rows and the columns and multiplying by $-\frac{1}{2}$.
3. Compute the eigenvalue decomposition $S = U\Lambda U^T$.
4. Let \tilde{X} be the matrix containing the first r columns of $\Lambda^{1/2}U^T$.

```
library(dimRed)
```

```
isomap_sr <- embed(cbind(X, Y, Z), "Isomap", knn = 10,  
                  ndim = 2)
```

```
## 2019-11-20 10:08:31: Isomap START
```


```
## 2019-11-20 10:08:31: constructing knn graph
```

```
## 2019-11-20 10:08:31: calculating geodesic distances
```

```
## 2019-11-20 10:08:31: Classical Scaling
```

```
isomap_sr@data@data %>%  
  plot(col = as.character(colours), pch = 19)
```

Swiss roll iii



Example i

```
isomap_res <- embed(data, "Isomap", knn = 10,  
                    ndim = 2)
```

```
## 2019-11-20 10:08:32: Isomap START
```

```
## 2019-11-20 10:08:32: constructing knn graph
```


```
## 2019-11-20 10:08:52: calculating geodesic distances
```

```
## 2019-11-20 10:09:05: Classical Scaling
```

Example ii

```
isomap_res@data %>%  
  as.data.frame() %>%  
  ggplot(aes(iso.1, iso.2)) +  
  geom_point(alpha = 0.5) +  
  theme_minimal()
```


Example iii



Example iv

ISO1=-6638



ISO1=-4258



ISO1=-3369



ISO1=-2722



ISO1=-2114



ISO1=-1504



ISO1=-818



ISO1=-150



ISO1=548



ISO1=1322



ISO1=2017



ISO1=2762



ISO1=3589



ISO1=4558




ISO1=6831



Intuition i

- The reason why Isomap works is because the shortest path distance approximates the **geodesic** distance on the manifold
 - “Train tracks distance”
- If we embed the weighted graph in \mathbb{R}^p , with each nodes at its corresponding data point, and each edge having length equal to the Euclidean distance, we can see the graph as a scaffold of the manifold.
- As we increase the sample size, the scaffold “converges” to the actual manifold.

Intuition ii



Tenenbaum *et al.* Science (2000)

Further examples i

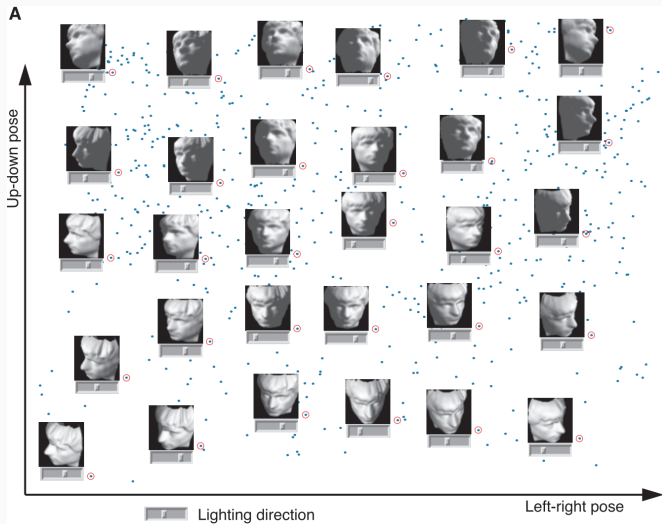


Figure 1

Further examples ii




Figure 2

Comments

- Advantages:
 - Simple extension of MDS
 - Preserves distance relationship on the manifold
- Disadvantages:
 - Computing the shortest path distance can be expensive with many data points
 - Doesn't work well with all manifolds (e.g. it fails when the underlying manifold has holes or many folds)