

Multidimensional Scaling

Max Turgeon

STAT 4690—Applied Multivariate Analysis

Recap: PCA

- We discussed several interpretations of PCA.
 - **Pearson:** PCA gives the best linear approximation to the data (at a fixed dimension).
- We also used PCA to visualized multivariate data:
 - Fit PCA
 - Plot PC1 against PC2.

Multidimensional scaling

- **Multidimensional scaling** is a method that looks at these two goals explicitly.
 - It has PCA as a special case.
 - But it is much more general.
- The input of MDS is a **dissimilarity matrix** Δ , and it aims to represent the data in a lower-dimensional space such that the resulting dissimilarities $\tilde{\Delta}$ are as close as possible to the original dissimilarities.
 - $\Delta \approx \tilde{\Delta}$.

Example of dissimilarities

- Dissimilarities measure how *different* two observations are.
 - Larger dissimilarity, more different.
- Therefore, any distance measure can be used as a dissimilarity measure.
 - Euclidean distance in \mathbb{R}^p .
 - Mahalanobis distance.
 - Driving distance between cities.
 - Graph-based distance.
- Any *similarity* measure can be turned into a dissimilarity measure using a monotone decreasing transformation.
 - E.g. $r_{ij} \implies 1 - r_{ij}^2$

Two types of MDS

- **Metric MDS**

- The embedding in the lower dimensional space uses the same dissimilarity measure as in the original space.

- **Nonmetric MDS**

- The embedding in the lower dimensional space only uses the rank information from the original space.

Metric MDS–Algorithm

- Input: An $n \times n$ matrix Δ of dissimilarities.
- Output: An $n \times r$ matrix \tilde{X} , with $r < p$.

Algorithm

1. Create the matrix D containing the square of the entries in Δ .
2. Create S by centering both the rows and the columns and multiplying by $-\frac{1}{2}$.
3. Compute the eigenvalue decomposition $S = U\Lambda U^T$.
4. Let \tilde{X} be the matrix containing the first r columns of $\Lambda^{1/2}U^T$.

Example i

```
Delta <- dist(swiss)
D <- Delta^2

# Center columns
B <- scale(D, center = TRUE, scale = FALSE)


# Center rows
B <- t(scale(t(B), center = TRUE, scale = FALSE))
B <- -0.5 * B
```

Example ii

```
decomp <- eigen(B)
Lambda <- diag(pmax(decomp$values, 0))
X_tilde <- decomp$vectors %*% Lambda^0.5

plot(X_tilde)
```


Example iii



Example iv

```
mds <- cmdscale(Delta, k = 2)

all.equal(X_tilde[,1:2], mds,
          check.attributes = FALSE)

## [1] TRUE
```

Example v

```
library(tidyverse)
# Let's add annotations
dimnames(X_tilde) <- list(rownames(swiss),
                           paste0("MDS", seq_len(ncol(X_tilde))))
X_tilde <- as.data.frame(X_tilde) %>%
  rownames_to_column("District")
```

Example vi

```
X_tilde <- X_tilde %>%  
  mutate(Canton = case_when(  
    District %in% c("Courtelary", "Moutier",  
                    "Neuveville") ~ "Bern",  
    District %in% c("Broye", "Glane", "Gruyere",  
                    "Sarine", "Veveyse") ~ "Fribourg",  
    District %in% c("Conthey", "Entremont", "Herens",  
                    "Martigwy", "Monthey",  
                    "St Maurice", "Sierre",  
                    "Sion") ~ "Valais"))
```


Example vii

```
X_tilde <- X_tilde %>%  
  mutate(Canton = case_when(!is.na(Canton) ~ Canton,  
    District %in% c("Boudry", "La Chauxdfnd",  
      "Le Locle", "Neuchatel",  
      "ValdeTravers",  
      "Val de Ruz") ~ "Neuchatel",  
    District %in% c("V. De Geneve", "Rive Droite",  
      "Rive Gauche") ~ "Geneva",  
    District %in% c("Delemont", "Franches-Mnt",  
      "Porrentruy") ~ "Jura",  
    TRUE ~ "Vaud"))
```

Example viii

```
library(ggrepel)
X_tilde %>%
  ggplot(aes(MDS1, MDS2)) +
  geom_point(aes(colour = Canton)) +
  geom_label_repel(aes(label = District)) +
  theme_minimal() +
  theme(legend.position = "top")
```

Example ix



Geographical distribution of the languages of Switzerland (2000)

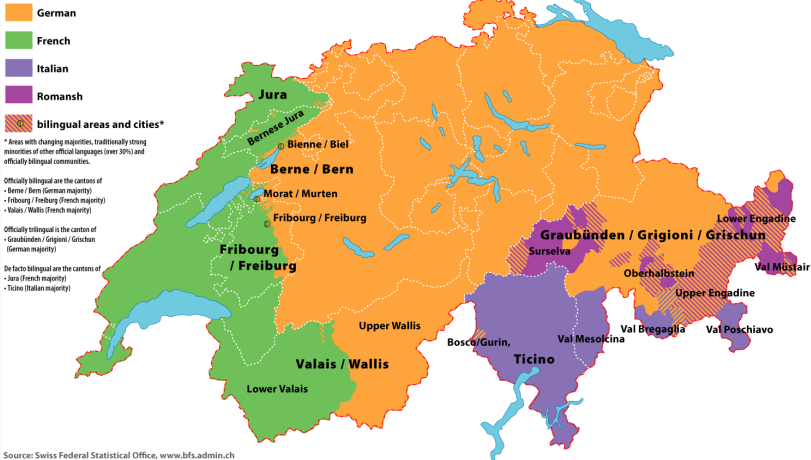


Figure 1

Another example i

```
library(psych)  
cities[1:5, 1:5]
```

##	ATL	BOS	ORD	DCA	DEN
## ATL	0	934	585	542	1209
## BOS	934	0	853	392	1769
## ORD	585	853	0	598	918
## DCA	542	392	598	0	1493
## DEN	1209	1769	918	1493	0

Another example ii


```
mds <- cmdscale(cities, k = 2)
colnames(mds) <- c("MDS1", "MDS2")

mds <- mds %>%
  as.data.frame %>%
  rownames_to_column("Cities")
```

Another example iii

```
mds %>%  
  ggplot(aes(MDS1, MDS2)) +  
  geom_point() +  
  geom_label_repel(aes(label = Cities)) +  
  theme_minimal()
```


Another example iv



Another example v

```
mds %>%  
  mutate(MDS1 = -MDS1, MDS2 = -MDS2) %>%  
  ggplot(aes(MDS1, MDS2)) +  
  geom_point() +  
  geom_label_repel(aes(label = Cities)) +  
  theme_minimal()
```

Another example vi



Why does it work? i

- The algorithm may seem like black magic...
 - Double centering?
 - Eigenvectors of distances?
- Let's try to justify it.
- Let $\mathbf{Y}_1, \dots, \mathbf{Y}_n$ be a set of points in \mathbb{R}^p .
- Recall that in \mathbb{R}^p , the Euclidean distance and the scalar product are related as follows:

Why does it work? ii

$$\begin{aligned}d(\mathbf{Y}_i, \mathbf{Y}_j)^2 &= \langle \mathbf{Y}_i - \mathbf{Y}_j, \mathbf{Y}_i - \mathbf{Y}_j \rangle \\&= (\mathbf{Y}_i - \mathbf{Y}_j)^T (\mathbf{Y}_i - \mathbf{Y}_j) \\&= \mathbf{Y}_i^T \mathbf{Y}_i - 2\mathbf{Y}_i^T \mathbf{Y}_j + \mathbf{Y}_j^T \mathbf{Y}_j.\end{aligned}$$

- In other words, the scalar product between \mathbf{Y}_i and \mathbf{Y}_j is given by

$$\mathbf{Y}_i^T \mathbf{Y}_j = -\frac{1}{2} \left(d(\mathbf{Y}_i, \mathbf{Y}_j)^2 - \mathbf{Y}_i^T \mathbf{Y}_i - \mathbf{Y}_j^T \mathbf{Y}_j \right).$$

Why does it work? iii

- Let S be the matrix whose (i, j) -th entry is $\mathbf{Y}_i^T \mathbf{Y}_j$, and note that D is the matrix whose (i, j) -th entry is $d(\mathbf{Y}_i, \mathbf{Y}_j)^2$.
- Now, assume that the dataset $\mathbf{Y}_1, \dots, \mathbf{Y}_n$ has sample mean $\bar{\mathbf{Y}} = 0$ (i.e. it is centred). The average of the i -th row of D is

Why does it work? iv

$$\begin{aligned}\frac{1}{n} \sum_{j=1}^n d(\mathbf{Y}_i, \mathbf{Y}_j)^2 &= \frac{1}{n} \sum_{j=1}^n (\mathbf{Y}_i^T \mathbf{Y}_i - 2\mathbf{Y}_i^T \mathbf{Y}_j + \mathbf{Y}_j^T \mathbf{Y}_j) \\&= \mathbf{Y}_i^T \mathbf{Y}_i - \frac{2}{n} \sum_{j=1}^n \mathbf{Y}_i^T \mathbf{Y}_j + \frac{1}{n} \sum_{j=1}^n \mathbf{Y}_j^T \mathbf{Y}_j \\&= \mathbf{Y}_i^T \mathbf{Y}_i - 2\mathbf{Y}_i^T \bar{\mathbf{Y}} + \frac{1}{n} \sum_{j=1}^n \mathbf{Y}_j^T \mathbf{Y}_j \\&= S_{ii} + \frac{1}{n} \sum_{j=1}^n S_{jj}.\end{aligned}$$

Why does it work? v

- Similarly, the average of the j -th column of D is given by

$$\frac{1}{n} \sum_{i=1}^n d(\mathbf{Y}_i, \mathbf{Y}_j)^2 = \frac{1}{n} \sum_{i=1}^n S_{ii} + S_{jj}.$$

- We can then deduce that the mean of **all** the entries of D is given by

$$\frac{1}{n^2} \sum_{i=1}^n \sum_{j=1}^n d(\mathbf{Y}_i, \mathbf{Y}_j)^2 = \frac{1}{n} \sum_{i=1}^n S_{ii} + \frac{1}{n} \sum_{j=1}^n S_{jj}.$$

Why does it work? vi

- Putting all of this together, we now have that

$$\begin{aligned}\mathbf{Y}_i^T \mathbf{Y}_i + \mathbf{Y}_j^T \mathbf{Y}_j &= \frac{1}{n} \sum_{j=1}^n d(\mathbf{Y}_i, \mathbf{Y}_j)^2 \\ &\quad + \frac{1}{n} \sum_{i=1}^n d(\mathbf{Y}_i, \mathbf{Y}_j)^2 \\ &\quad - \frac{1}{n^2} \sum_{i=1}^n \sum_{j=1}^n d(\mathbf{Y}_i, \mathbf{Y}_j)^2.\end{aligned}$$

Why does it work? vii

- **In other words**, we can recover the scalar products from the square distances through double centering and scaling.
- Moreover, since we assumed the data was centred, the matrix S is proportional to the sample covariance matrix.
 - In this context, up to a constant, MDS and PCA give the same results.
- **Note:** This idea that double centering allows us to go from dissimilarities to scalar products will come back again in the next lecture on kernel methods.

Further comments

- In PCA, we performed an eigendecomposition of the sample covariance matrix.
 - This is a $p \times p$ matrix.
- In MDS, we performed an eigendecomposition of the doubly centred and scaled matrix of squared distances.
 - This is an $n \times n$ matrix.
- If our dissimilarities are computed using the Euclidean distance, both methods will give the same answer.
 - **BUT**: the smallest matrix will be faster to compute and faster to decompose.
 - $n > p \Rightarrow$ PCA; $n < p \Rightarrow$ MDS

Stress function

- Nonmetric MDS approaches the problem a bit differently.
- We still have the same output Δ of dissimilarities, but we also have an objective function called the **stress function**.
- Recall that our goal is to represent the data in a lower-dimensional space such that the resulting dissimilarities $\tilde{\Delta}$ are as close as possible to the original dissimilarities.
 - $\Delta_{ij} \approx \tilde{\Delta}_{ij}$, for all i, j .

Stress function ii

- The stress function is defined as

$$\text{Stress}(\tilde{\Delta}; r) = \sqrt{\frac{\sum_{i,j=1}^n w_{ij} (\Delta_{ij} - \tilde{\Delta}_{ij})^2}{c}},$$

where

- w_{ij} are nonnegative weights;
- c is a normalising constant.
- Note that the stress function depends on both the dimension r of the lower space and the distances $\tilde{\Delta}$.
- **Goal:** Find points in \mathbb{R}^r such that their similarities minimise the stress function.

Sammon's Nonlinear Mapping

- The stress function is

$$\text{Stress}(\tilde{\Delta}; r) = \frac{1}{c} \sum_{i=1, i < j}^n \frac{(\Delta_{ij} - \tilde{\Delta}_{ij})^2}{\Delta_{ij}},$$

where

$$c = \sum_{i=1, i < j}^n \Delta_{ij}.$$

- We don't make any assumption on the dissimilarities Δ , but we assume that $\tilde{\Delta}$ arises from the Euclidean distance in \mathbb{R}^r .
 - This makes the minimisation problem easier and amenable to Newton's method.

Example i

```
library(MASS)
```

```
Delta <- dist(swiss)
```

```
mds <- sammon(Delta, k = 2)
```


```
## Initial stress          : 0.01959
```

```
## stress after    0 iters: 0.01959
```

Example ii

```
plot(mds$points)
```


Example iii



Example iv

```
# Fit for different values of k
stresses <- sapply(seq(2, 10),
  function(k) {
    sammon(Delta, k = k,
      trace = FALSE)$stress
  })
plot(seq(2, 10), stresses, type = 'b')
```

Example v




Example vi

```
library(scatterplot3d)
mds <- sammon(Delta, k = 3)

## Initial stress          : 0.00243
## stress after 10 iters: 0.00095, magic = 0.500
## stress after 20 iters: 0.00094, magic = 0.500

scatterplot3d(mds$points,
              xlab = "MDS1", ylab = "MDS2",
              zlab = "MDS3")
```

Example vii



Kruskal's Nonmetric MDS

- Kruskal's approach is based on **ranks**.
- In other words: instead of finding points in \mathbb{R}^r with similar distances, his method tries to preserve the relative ordering of the dissimilarities.
 - The most dissimilar points in \mathbb{R}^p should be represented by the most dissimilar points in \mathbb{R}^r , but the actual magnitude is irrelevant.
- This is achieved by allowing a monotone transformation f of the dissimilarities. We thus get

$$\text{Stress}(\tilde{\Delta}; r) = \sqrt{\frac{\sum_{i=1, i < j}^n (\Delta_{ij} - f(\tilde{\Delta}_{ij}))^2}{\sum_{i=1, i < j}^n \Delta_{ij}}}.$$

Example (cont'd) i

```
mds_s <- sammon(Delta, k = 2)
```

```
## Initial stress          : 0.01959
```

```
## stress after    0 iters: 0.01959
```



```
mds_k <- isoMDS(Delta, k = 2)
```

Example (cont'd) ii

```
## initial value 5.463800
## iter 5 value 4.499103
## iter 5 value 4.495335
## iter 5 value 4.492669
## final value 4.492669
## converged
```

```
par(mfrow = c(1, 2))
plot(mds_s$points, main = "Sammon",
     xlab = "MDS1", ylab = "MDS2")
plot(mds_k$points, main = "Kruskal",
     xlab = "MDS1", ylab = "MDS2")
```


Example (cont'd) iii



Example (cont'd) iv

```
# Sammon and Kruskal have different  
# optimal k  
stresses <- sapply(seq(2, 10),  
                   function(k) {  
                     isoMDS(Delta, k = k,  
                             trace = FALSE)$stress  
                   })  
plot(seq(2, 10), stresses, type = 'b')
```

Example (cont'd) v





Example (cont'd) vi

```
mds_opt_s <- sammon(Delta, k = 3,  
                    trace = FALSE)  
mds_opt_k <- isoMDS(Delta, k = 6,  
                    trace = FALSE)  
  
# Let's cluster in the MDS space  
cluster_s <- kmeans(mds_opt_s$points, centers = 2)  
cluster_k <- kmeans(mds_opt_k$points, centers = 2)
```

Example (cont'd) vii

```
par(mfrow = c(1, 2))  
plot(mds_s$points, main = "Sammon",  
      xlab = "MDS1", ylab = "MDS2",  
      col = cluster_s$cluster)  
plot(mds_k$points, main = "Kruskal",  
      xlab = "MDS1", ylab = "MDS2",  
      col = cluster_k$cluster)
```



Example (cont'd) viii



Example (cont'd) ix

```
# More interestingly, you can use MDS to  
# cluster data where you only have distances  
stresses <- sapply(seq(2, 6),  
                  function(k) {  
                    isoMDS(as.matrix(cities),  
                           k = k,  
                           trace = FALSE)$stress  
                  })  
plot(seq(2, 6), stresses, type = 'b')
```


Example (cont'd) x



Example (cont'd) xi

```
mds_cities <- isoMDS(as.matrix(cities), k = 6,  
                    trace = FALSE)  
cluster_cities <- kmeans(mds_cities$points,  
                        centers = 2)  
  
plot(mds_cities$points, main = "Kruskal",  
     xlab = "MDS1", ylab = "MDS2",  
     type = 'n')  
text(mds_cities$points, colnames(cities),  
     col = cluster_cities$cluster)
```

Example (cont'd) xii



Summary

- Multidimensional scaling is mainly a method for visualising multivariate data.
- It works by finding points in a lower dimensional space with similar dissimilarities than the one on the original space.
- It only requires a matrix of dissimilarities
 - Therefore, it allows us to visualise data with limited information.
- MDS is an example of a **nonlinear dimension reduction** method.