

# Factor Analysis

---

Max Turgeon

STAT 4690—Applied Multivariate Analysis

# Latent variable models

- With PCA, we saw how we could reduce the dimension of data using the eigenvectors of the sample covariance matrix.
- Conversely, we could construe PCA has a generative model, where the principal components give rise to the observed data.
- **Latent Variable Models** formalise this idea:
  - Latent (i.e. unobserved) variables  $\mathbf{F}$  give rise to observed data  $\mathbf{Y}$  through a *specified* model.

# Factor Analysis i

- **Factor Analysis** is a special kind of latent variable model.
- Let  $\mathbf{Y}$  be a  $p$ -dimensional vector with mean  $\mu$  and covariance matrix  $\Sigma$ .
- Let  $\mathbf{F}$  be a  $m$ -dimensional *latent* vector.
- The *orthogonal factor model* is given by

$$\mathbf{Y} - \mu = L\mathbf{F} + \mathbf{E},$$

where  $L$  is a  $p \times m$  *matrix of factor loadings*, and  $\mathbf{E}$  is a  $p$ -dimensional vector of *errors*.

## Factor Analysis ii

- $F$  are also called *common factors*;  $E$  are also called *specific factors*.
- **Note:** This is essentially a multivariate regression model, but where the covariates are unobserved.

# Assumptions i

- The model above is generally not identifiable, since there are too many parameters.
- We therefore need to impose further restrictions:
  - $E(\mathbf{F}) = 0$
  - $\text{Cov}(\mathbf{F}) = I$
  - $E(\mathbf{E}) = 0$
  - $\text{Cov}(\mathbf{E}) = \Psi = \begin{pmatrix} \psi_1 & 0 & \cdots & 0 \\ 0 & \psi_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \psi_p \end{pmatrix}$
  - $\text{Cov}(\mathbf{F}, \mathbf{E}) = 0$

## Assumptions ii

- In other words:
- Both common and specific factors have mean zero;
- They are uncorrelated;
- The common factors are mutually uncorrelated and standardised;
- The specific factors each affect only one observed variable.

# Structured Covariance i

- As a consequence of these assumptions, we can derive an assumption on the structure of  $\Sigma = \text{Cov}(\mathbf{Y})$ :

$$\begin{aligned}\Sigma &= E\left((\mathbf{Y} - \mu)(\mathbf{Y} - \mu)^T\right) \\ &= E\left((L\mathbf{F} + \mathbf{E})(L\mathbf{F} + \mathbf{E})^T\right) \\ &= LE(\mathbf{F}\mathbf{F}^T)L + E(\mathbf{E}\mathbf{F}^T)L^T + LE(\mathbf{F}\mathbf{E}^T) + E(\mathbf{E}\mathbf{E}^T) \\ &= LIL^T + 0L^T + L0 + \Psi \\ &= LL^T + \Psi.\end{aligned}$$

## Structured Covariance ii

- Similarly, we can show that

$$\text{Cov}(\mathbf{Y}, \mathbf{F}) = L.$$

- If we write  $\ell_{ij}$  for the  $(i, j)$ -th element of  $L$ , we see that

$$\text{Var}(Y_i) = \sum_{k=1}^m \ell_{ik}^2 + \psi_i.$$

- Crucially, these assumptions are **testable**. In other words, we can check whether they are reasonable for our data.



## Example i

- Let's look at an example where there is no solution.
- Assume  $p = 3$ ,  $m = 1$ , with

$$\Sigma = \begin{pmatrix} 1 & 0.9 & 0.7 \\ 0.9 & 1 & 0.4 \\ 0.7 & 0.4 & 1 \end{pmatrix}.$$

- From our assumptions on the covariance structure, we derive a system of equations

$$\begin{aligned} 1 &= \ell_{11}^2 + \psi_1 & 0.9 &= \ell_{11}\ell_{21} & 0.7 &= \ell_{11}\ell_{31} \\ & & 1 &= \ell_{22}^2 + \psi_2 & 0.4 &= \ell_{21}\ell_{31} \\ & & & & 1 &= \ell_{33}^2 + \psi_3 \end{aligned}$$

## Example ii

- From  $0.7 = \ell_{11}\ell_{31}$  and  $0.4 = \ell_{21}\ell_{31}$ , we get

$$\ell_{21} = \frac{0.4}{0.7}\ell_{11}.$$

- But since  $0.9 = \ell_{11}\ell_{21}$ , we can conclude that

$$\ell_{11} = \pm 1.255.$$

- However, since the first component  $Y_1$  has unit variance,  $\ell_{11} = \text{Corr}(Y_1, F_1)$ , and therefore the correlation is out of bounds.
- Similarly, we get

$$\psi_1 = 1 - \ell_{11}^2 = 1 - 1.575 = -0.575.$$

## Example iii

- But since  $\psi_1$  is the variance of the first error term, we once again get a non-sensical solution.

# Factor Rotation i

- Even with our assumptions above, our model is still not uniquely identified.
- Let  $T$  be an  $m \times m$  orthogonal matrix. We have

$$\begin{aligned}\mathbf{Y} - \mu &= L\mathbf{F} + \mathbf{E} \\ &= LTT^T\mathbf{F} + \mathbf{E} \\ &= \tilde{L}\tilde{\mathbf{F}} + \mathbf{E},\end{aligned}$$

where  $\tilde{L} = LT$  and  $\tilde{\mathbf{F}} = T^T\mathbf{F}$ .

## Factor Rotation ii

- Both models lead to the same covariance matrix:

$$\Sigma = LL^T + \Psi = LTT^TL^T + \Psi = \tilde{L}\tilde{L}^T + \Psi.$$

- As we will see, this will turn out to be a blessing in disguise:
  - We will impose a uniqueness condition to get one solution.
  - Then we will rotate our solution using  $T$  to improve interpretation.

# Estimation—Principal Component Method i

- Recall the spectral decomposition of the covariance matrix:

$$\Sigma = \sum_{i=1}^p \lambda_i w_i w_i^T,$$

with  $\lambda_1 \geq \dots \geq \lambda_p$ .

- If we let  $W$  be the matrix whose  $i$ -th column is  $\sqrt{\lambda_i} w_i$ , we can rewrite the spectral decomposition as

$$\Sigma = WW^T.$$

- In other words, if we let  $m = p$  and  $\Psi = 0$ , we see that we recover the orthogonal factor model with  $L = W$ .

## Estimation—Principal Component Method ii

- Of course, this is not very satisfactory, as the dimension of the common factors is the same as that of the original data.
- Instead, we select  $m < p$  using one of the methods we discussed with PCA and we approximate

$$\Sigma \approx \sum_{i=1}^m \lambda_i w_i w_i^T.$$

- If we let  $L$  be the  $p \times m$  matrix whose  $i$ -th column is  $\sqrt{\lambda_i} w_i$ , we can estimate  $\Psi$  as follows:

$$\psi_i = \sigma_{ii} - \sum_{j=1}^m \ell_{ij}^2.$$

## Algorithm

1. Let  $\hat{\lambda}_1 \cdots > \hat{\lambda}_p$  and  $\hat{w}_1, \dots, \hat{w}_p$  be the eigenvalues and eigenvectors of the covariance matrix  $S_n$ .
2. Select  $m$  using one of the PCA criteria.
3. Estimate  $\hat{L}$  with the  $p \times m$  matrix whose  $i$ -th column is  $\sqrt{\hat{\lambda}_i} \hat{w}_i$ .
4. Estimate  $\hat{\Psi}$  with the diagonal elements of  $S_n - \hat{L}\hat{L}^T$ .



## Example i

```
library(psych)
```

```
dim(bfi)
```

```
## [1] 2800 28
```

```
names(bfi)
```

## Example ii

##	[1]	"A1"	"A2"	"A3"	"A4"
##	[6]	"C1"	"C2"	"C3"	"C4"
##	[11]	"E1"	"E2"	"E3"	"E4"
##	[16]	"N1"	"N2"	"N3"	"N4"
##	[21]	"O1"	"O2"	"O3"	"O4"
##	[26]	"gender"	"education"	"age"	

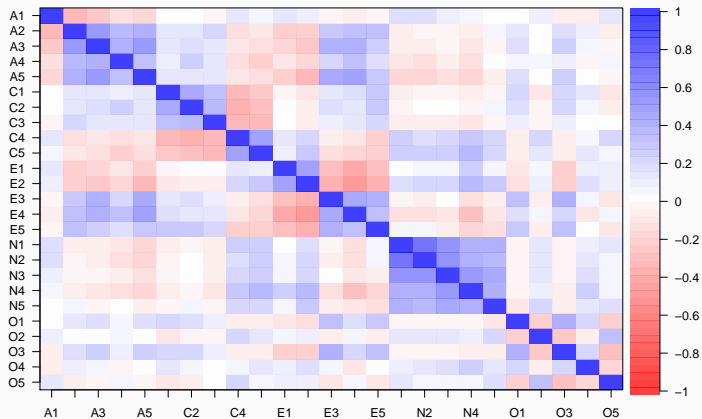
```
library(tidyverse)
```

```
data <- bfi %>%  
  select(-gender, -education, -age) %>%  
  filter(complete.cases(.))
```

## Example iii

```
cor.plot(cor(data))
```

## Example iv



## Example v

```
decomp <- prcomp(data)
summary(decomp)$importance[,1:3]
```

##	PC1	PC2	PC3
## Standard deviation	3.291635	2.451538	2.030393
## Proportion of Variance	0.215650	0.119620	0.082050
## Cumulative Proportion	0.215650	0.335270	0.417320

## Example vi

```
cum_prop <- decomp %>%  
  summary %>%  
  .[["importance"]] %>%  
  .["Cumulative Proportion",]
```

```
which(cum_prop > 0.8)[1]
```

```
## PC14
```

```
## 14
```

## Example vii

```
Lhat <- decomp$rotation[,1:14] %*%  
  diag(decomp$sdev[1:14])  
Psi_hat <- diag(cov(data) - tcrossprod(Lhat))  
  
# Sum squared error  
sum((cov(data) - tcrossprod(Lhat) - diag(Psi_hat))^2)  
  
## [1] 3.645694  
  
# Compare to the total variance  
sum(diag(cov(data)))
```

## Example viii

```
## [1] 50.24287
```

```
# Our FA model explains:
```

```
sum(colSums(Lhat^2)/sum(diag(cov(data))))
```

```
## [1] 0.8160488
```



# Comments

- In our example above, we saw that 14 factors explained 82% of the total variance.
- Two common default values for  $m$  in statistical softwares:
  - Number of positive eigenvalues of the sample covariance matrix.
  - Number of eigenvalues greater than one for the sample correlation matrix.
- In our example, the first criterion would lead to  $m = p$ , which is not very helpful

## Example (cont'd) i

```
(m <- sum(eigen(cor(data))$values > 1))
```

```
## [1] 6
```

```
Lhat <- decomp$rotation[,seq_len(m)] %*%  
  diag(decomp$sdev[seq_len(m)])
```

```
Psi_hat <- diag(cov(data) - tcrossprod(Lhat))
```

```
# Sum squared error
```

```
sum((cov(data) - tcrossprod(Lhat) - diag(Psi_hat))^2)
```

## Example (cont'd) ii

```
## [1] 6.999035
```

```
# Compare to the total variance  
sum(diag(cov(data)))
```

```
## [1] 50.24287
```

```
# Our FA model explains:  
sum(colSums(Lhat2)/sum(diag(cov(data))))
```

```
## [1] 0.5910586
```

## Example (cont'd) iii

*# We can also visualize the fit*


```
Sn <- cov(data)
```

```
Sn_fit <- tcrossprod(Lhat) - diag(Psi_hat)
```

```
library(lattice)
```

```
levelplot(Sn - Sn_fit)
```

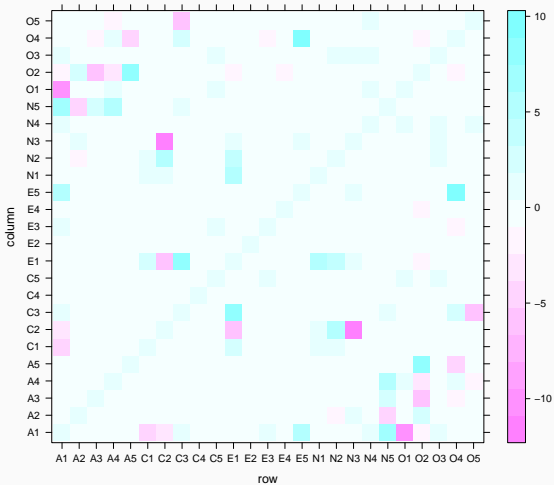
## Example (cont'd) iv



## Example (cont'd) v

```
# Or if you prefer % difference  
levelplot((Sn - Sn_fit)/Sn)
```

## Example (cont'd) vi



# Estimation—Maximum Likelihood Method i

- This estimation method assumes that both common factors  $\mathbf{F}$  and specific factors  $\mathbf{E}$  follow a multivariate normal distribution
  - $\mathbf{F} \sim N_m(0, I)$
  - $\mathbf{E} \sim N_p(0, \Psi)$
- From this assumption, it follows that  $\mathbf{Y}$  also follows a multivariate normal distribution
  - $\mathbf{Y} \sim N_p(\mu, LL^T + \Psi)$
- Therefore, we can write down the likelihood in terms of both  $L$  and  $\Psi$ .



## Estimation—Maximum Likelihood Method ii

- However, because of the factor rotation problem, we need to impose a constraint in order to obtain a unique solution:
  - $L^T \Psi^{-1} L = \Delta$  is diagonal.
- Given this assumption, the maximum likelihood estimates  $\hat{L}$  and  $\hat{\Psi}$  can be found using an iterative algorithm.
- We will not go into the details of the algorithm (but see Johnson & Wichern, Supplement 9A if interested).
  - Instead, we will rely on the R function `stats::factanal`.

## Example (cont'd) i

```
# CAREFUL: uses correlation matrix  
fa_decomp <- factanal(data, factors = m,  
                      rotation = 'none')  
  
# Uniquenesses are the diagonal elements  
# of the matrix Psi  
Psi_mle <- fa_decomp$uniquenesses  
Lmle <- fa_decomp$loadings
```

## Example (cont'd) ii

*# We get an estimate of the correlation*

```
R_mle <- tcrossprod(Lmle) - diag(Psi_mle)
```

```
sum((cor(data) - R_mle)^2)
```

```
## [1] 31.97525
```

*# Our FA model explains:*

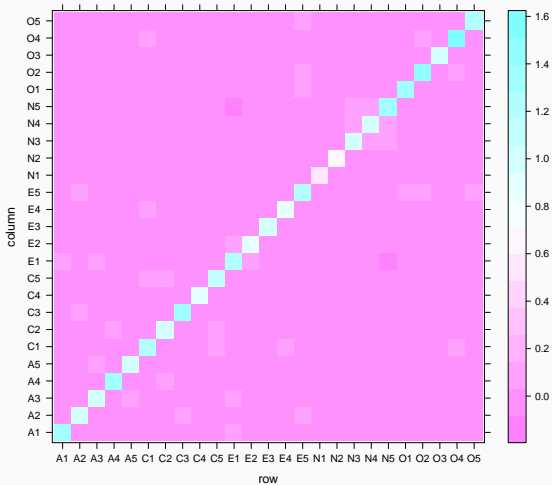
```
sum(colSums(Lmle^2)/ncol(data))
```

```
## [1] 0.4498342
```

## Example (cont'd) iii

```
levelplot(cor(data) ~ R_mle)
```

## Example (cont'd) iv



## Example (cont'd) v

```
# To factor the covariance matrix
```

```
# Use psych::fa
```

```
fa_decomp <- psych::fa(data, nfactors = m,  
                        rotate = "none",  
                        covar = TRUE,  
                        fm = "ml")
```

```
# Extract estimates
```

```
Psi_mle <- fa_decomp$uniquenesses
```

```
Lmle <- fa_decomp$loadings
```

## Example (cont'd) vi

*# We get an estimate of the covariance*

```
Sn_mle <- tcrossprod(Lmle) - diag(Psi_mle)
```

```
sum((Sn - Sn_mle)^2)
```

```
## [1] 128.025
```

*# Our FA model explains:*

```
sum(colSums(Lmle^2)/sum(diag(Sn)))
```

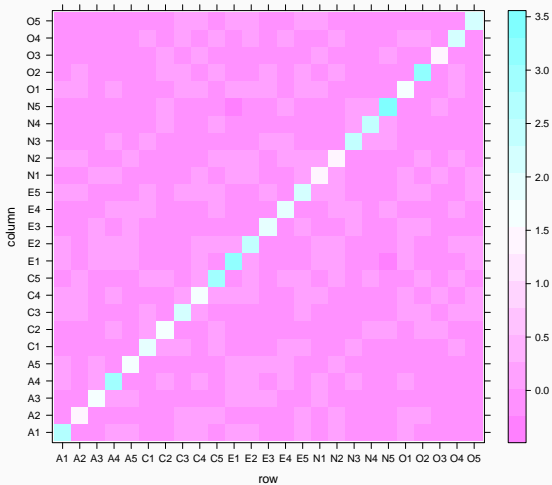
```
## [1] 0.459665
```

## Example (cont'd) vii

```
levelplot(Sn - Sn_mle)
```




## Example (cont'd) viii



## Example (cont'd) ix

```
# Compare MLE with PC estimate  
levelplot((Sn_fit - Sn_mle)/Sn)
```

## Example (cont'd) x



# Comments about estimation

- There are other methods of estimating the loadings and the “uniquenesses”
  - Ordinary Least Squares
  - Weighted OLS
  - Principal factor
- With so many choices of estimation methods, it can be hard to compare statistical softwares
  - And you also have to read the manual in order to know what is going on...
- You also always have the choice between factoring the covariance or the correlation matrix
  - Which one you choose depends on the commensurability of your variables (just like in PCA)
- Finally, it's always a good idea to compare the output of

# Factor Rotation Redux i

- As we saw earlier, any orthogonal matrix  $T$  gives rise to the same factor analysis model

$$\Sigma = LL^T + \Psi = LTT^TL^T + \Psi = \tilde{L}\tilde{L}^T + \Psi.$$

- In other words, we cannot choose  $T$  to maximise the goodness of fit.
  - We need another criterion
- Intuitively, to ease interpretation, we want each variable to have large loadings for one factor and negligible loadings for the other ones.

## Factor Rotation Redux ii

- [https://maxturgeon.ca/f19-stat4690/factor\\_rotation.gif](https://maxturgeon.ca/f19-stat4690/factor_rotation.gif)
- One common analytic criterion that formalises this idea is the **varimax criterion**.
  - Resulting loadings are called *varimax loadings*
- We have

$$\text{VARIMAX} \propto \sum_{j=1}^m \left( \begin{array}{c} \text{Variance of squares of scales loadings} \\ \text{for } j\text{-th factor} \end{array} \right)$$

- More precisely:
  - Let  $\tilde{\ell}_{ij}$  be the  $(i, j)$ -th entry of the matrix  $\tilde{L} = LT$ . In other words,  $\tilde{\ell}_{ij}$  depends on  $T$ .

## Factor Rotation Redux iii

- Let  $\tilde{h}_i^2 = \sum_{j=1}^m \tilde{\ell}_{ij}^2$ .
- Define the scaled loadings  $\tilde{\ell}_{ij}^* = \tilde{\ell}_{ij}/h_i$ .
- The varimax criterion  $V$  is given as

$$V = \frac{1}{p} \sum_{j=1}^m \left( \sum_{i=1}^p \tilde{\ell}_{ij}^{*4} - \frac{1}{p} \left( \sum_{i=1}^p \tilde{\ell}_{ij}^{*2} \right)^2 \right).$$



- In R, you can compute the rotated loadings using the `stats::varimax` function. Alternatively, the function `stats::factanal` can compute the rotation for you as part of the factor analysis (and so does `psych::fa`).

## Example (cont'd) i

```
# Let's start with m=2 for visualization  
fa_decomp <- factanal(data, factors = 2,  
                      rotation = 'none')  
  
initial_loadings <- fa_decomp$loadings  
varimax_loadings <- varimax(initial_loadings)
```



## Example (cont'd) ii



## Example (cont'd) iii

*# You can extract the matrix T*

```
varimax_loadings$rotmat
```

```
##           [,1]      [,2]
```

```
## [1,] 0.7810310 -0.6244923
```

```
## [2,] 0.6244923  0.7810310
```

*# We can also get the angle of rotation*

```
acos(varimax_loadings$rotmat[1,1])
```

```
## [1] 0.6744813
```


## Example (cont'd) iv

```
# In more dimensions
```

```
fa_decomp <- factanal(data, factors = m,  
                      rotation = 'varimax')
```

```
levelplot(unclass(fa_decomp$loadings),  
          xlab = "", ylab = "")
```

## Example (cont'd) v



# Comments

- As with estimation, there are many more rotation methods.
  - See for example the help page  
`?GPArotation::rotations`
- One particular class of rotations are called *oblique*
  - The matrix  $T$  is no longer constrained to be orthogonal.
- Factor rotation is especially useful with loadings obtained through MLE
  - Recall the constraint on  $L^T \Psi^{-1} L$  being diagonal
- Factor rotations are also sometimes used with PCA loadings.