

Review of Linear Algebra

Max Turgeon

STAT 4690—Applied Multivariate Analysis

Basic Matrix operations

Matrix algebra and R

- Matrix operations in R are *very* fast.
- This includes various class of operations:
 - Matrix addition, scalar multiplication, matrix multiplication, matrix-vector multiplication
 - Standard functions like determinant, rank, condition number, etc.
 - Matrix decompositions, e.g. eigenvalue, singular value, Cholesky, QR, etc.
 - Support for *sparse* matrices, i.e. matrices where a significant number of entries are exactly zero.

Matrix functions i

```
A <- matrix(c(1, 2, 3, 4), nrow = 2, ncol = 2)
```

```
A
```

```
##      [,1] [,2]
```

```
## [1,]    1    3
```

```
## [2,]    2    4
```

```
# Determinant
```

```
det(A)
```

```
## [1] -2
```

Matrix functions ii

Rank

```
library(Matrix)
```

```
rankMatrix(A)
```

```
## [1] 2
```

```
## attr(,"method")
```

```
## [1] "tolNorm2"
```

```
## attr(,"useGrad")
```

```
## [1] FALSE
```

```
## attr(,"tol")
```

```
## [1] 4.440892e-16
```

Matrix functions iii

Condition number

```
kappa(A)
```

```
## [1] 18.77778
```

How to compute the trace?

```
sum(diag(A))
```

```
## [1] 5
```

Matrix functions iv

```
# Transpose
```

```
t(A)
```

```
##      [,1] [,2]
```

```
## [1,]    1    2
```

```
## [2,]    3    4
```

```
# Inverse
```

```
solve(A)
```

Matrix functions v

```
##      [,1] [,2]  
## [1,]   -2  1.5  
## [2,]    1 -0.5
```

```
A %*% solve(A) # CHECK
```

```
##      [,1] [,2]  
## [1,]    1    0  
## [2,]    0    1
```


Matrix operations i

```
A <- matrix(c(1, 2, 3, 4), nrow = 2, ncol = 2)
B <- matrix(c(4, 3, 2, 1), nrow = 2, ncol = 2)
```

Addition

A + B

```
##      [,1] [,2]
## [1,]    5    5
## [2,]    5    5
```

Matrix operations ii

```
# Scalar multiplication
```

```
3*A
```

```
##      [,1] [,2]
```

```
## [1,]    3    9
```

```
## [2,]    6   12
```

```
# Matrix multiplication
```

```
A %*% B
```

Matrix operations iii

```
##      [,1] [,2]  
## [1,]   13   5  
## [2,]   20   8
```

Hadamard product aka entrywise multiplication

A * B

```
##      [,1] [,2]  
## [1,]    4   6  
## [2,]    6   4
```

Matrix operations iv

```
# Matrix-vector product
```

```
vect <- c(1, 2)
```

```
A %*% vect
```

```
##      [,1]
```

```
## [1,]    7
```

```
## [2,]   10
```

```
# BE CAREFUL: R recycles vectors
```

```
A * vect
```

Matrix operations v

```
##      [,1] [,2]  
## [1,]    1    3  
## [2,]    4    8
```

Eigenvalues and Eigenvectors

Eigenvalues

- Let \mathbf{A} be a square $n \times n$ matrix.
- The equation

$$\det(\mathbf{A} - \lambda I_n) = 0$$

is called the *characteristic equation* of \mathbf{A} .

- This is a polynomial equation of degree n , and its roots are called the *eigenvalues* of \mathbf{A} .

Example

Let

$$\mathbf{A} = \begin{pmatrix} 1 & 0.5 \\ 0.5 & 1 \end{pmatrix}.$$

Then we have

$$\begin{aligned} \det(\mathbf{A} - \lambda I_2) &= (1 - \lambda)^2 - 0.25 \\ &= (\lambda - 1.5)(\lambda - 0.5) \end{aligned}$$

Therefore, \mathbf{A} has two (real) eigenvalues, namely

$$\lambda_1 = 1.5, \lambda_2 = 0.5.$$

A few properties

Let $\lambda_1, \dots, \lambda_n$ be the eigenvalues of \mathbf{A} (with multiplicities).

1. $\text{tr}(\mathbf{A}) = \sum_{i=1}^n \lambda_i$;
2. $\det(\mathbf{A}) = \prod_{i=1}^n \lambda_i$;
3. The eigenvalues of \mathbf{A}^k are $\lambda_1^k, \dots, \lambda_n^k$, for k a nonnegative integer;
4. If \mathbf{A} is invertible, then the eigenvalues of \mathbf{A}^{-1} are $\lambda_1^{-1}, \dots, \lambda_n^{-1}$.

Eigenvectors

- If λ is an eigenvalue of \mathbf{A} , then (by definition) we have $\det(\mathbf{A} - \lambda I_n) = 0$.
- In other words, the following equivalent statements hold:
 - The matrix $\mathbf{A} - \lambda I_n$ is singular;
 - The kernel space of $\mathbf{A} - \lambda I_n$ is nontrivial (i.e. not equal to the zero vector);
 - The system of equations $(\mathbf{A} - \lambda I_n)v = 0$ has a nontrivial solution;
 - There exists a nonzero vector v such that

$$\mathbf{A}v = \lambda v.$$

- Such a vector is called an *eigenvector* of \mathbf{A} .

Example (cont'd) i

Recall that we had

$$\mathbf{A} = \begin{pmatrix} 1 & 0.5 \\ 0.5 & 1 \end{pmatrix},$$

and we determined that 0.5 was an eigenvalue of \mathbf{A} .

We therefore have

$$\mathbf{A} - 0.5I_2 = \begin{pmatrix} 0.5 & 0.5 \\ 0.5 & 0.5 \end{pmatrix}.$$

Example (cont'd) ii

As we can see, any vector v of the form $(x, -x)$ satisfies

$$(\mathbf{A} - 0.5I_2)v = (0, 0).$$

In other words, we not only get a single eigenvector, but a whole subspace of \mathbb{R}^2 . By convention, we usually select as a representative a vector of norm 1, e.g.

$$v = \left(\frac{1}{\sqrt{2}}, \frac{-1}{\sqrt{2}} \right).$$

Example (cont'd) iii

Alternatively, instead of finding the eigenvector by inspection, we can use the reduced row-echelon form of $\mathbf{A} - 0.5I_2$, which is given by

$$\mathbf{A} = \begin{pmatrix} 1 & 1 \\ 0 & 0 \end{pmatrix}.$$

Therefore, the solutions to $(\mathbf{A} - 0.5I_2)v$, with $v = (x, y)$ are given by a single equation, namely $y + x = 0$, or $y = -x$.

Eigenvalues and eigenvectors in R i

```
A <- matrix(c(1, 0.5, 0.5, 1), nrow = 2)
```

```
result <- eigen(A)
```

```
names(result)
```

```
## [1] "values" "vectors"
```

```
result$values
```

```
## [1] 1.5 0.5
```

Eigenvalues and eigenvectors in R ii

```
result$eigenvectors
```

```
##           [,1]      [,2]  
## [1,] 0.7071068 -0.7071068  
## [2,] 0.7071068  0.7071068
```

```
1/sqrt(2)
```

```
## [1] 0.7071068
```

Symmetric matrices i

- A matrix \mathbf{A} is called *symmetric* if $\mathbf{A}^T = \mathbf{A}$.
- **Proposition 1:** If \mathbf{A} is (real) symmetric, then its eigenvalues are real.

Proof: Let λ be an eigenvalue of \mathbf{A} , and let $v \neq 0$ be an eigenvector corresponding to this eigenvalue. Then we have

Symmetric matrices ii

$$\begin{aligned}\lambda \bar{v}^T v &= \bar{v}^T (\lambda v) \\ &= \bar{v}^T (\mathbf{A} v) \\ &= (\mathbf{A}^T \bar{v})^T v \\ &= (\mathbf{A} \bar{v})^T v && (\mathbf{A} \text{ is symmetric}) \\ &= (\overline{\mathbf{A} v})^T v && (\mathbf{A} \text{ is real}) \\ &= \bar{\lambda} \bar{v}^T v.\end{aligned}$$

Since we have $v \neq 0$, we conclude that $\lambda = \bar{\lambda}$, i.e. λ is real.



- **Proposition 2:** If A is (real) symmetric, then eigenvectors corresponding to distinct eigenvalues are orthogonal.

Proof: Let λ_1, λ_2 be distinct eigenvalues, and let $v_1 \neq 0, v_2 \neq 0$ be corresponding eigenvectors. Then we have

$$\begin{aligned}\lambda_1 v_1^T v_2 &= (\mathbf{A}v_1)^T v_2 \\ &= v_1^T \mathbf{A}^T v_2 \\ &= v_1^T \mathbf{A} v_2 \quad (\mathbf{A} \text{ is symmetric}) \\ &= v_1^T (\lambda_2 v_2) \\ &= \lambda_2 v_1^T v_2.\end{aligned}$$

Since $\lambda_1 \neq \lambda_2$, we conclude that $v_1^T v_2 = 0$, i.e. v_1 and v_2 are orthogonal. □

Spectral Decomposition i

- Putting these two propositions together, we get the *Spectral Decomposition* for symmetric matrices.
- **Theorem:** Let \mathbf{A} be an $n \times n$ symmetric matrix, and let $\lambda_1 \geq \dots \geq \lambda_n$ be its eigenvalues (with multiplicity).

Then there exist vectors v_1, \dots, v_n such that

1. $\mathbf{A}v_i = \lambda_i v_i$, i.e. v_i is an eigenvector, for all i ;
2. If $i \neq j$, then $v_i^T v_j = 0$, i.e. they are orthogonal;
3. For all i , we have $v_i^T v_i = 1$, i.e. they have unit norm;
4. We can write $\mathbf{A} = \sum_{i=1}^n \lambda_i v_i v_i^T$.

Sketch of a proof:

Spectral Decomposition ii

1. We are saying that we can find n eigenvectors. This means that if an eigenvalue λ has multiplicity m (as a root of the characteristic polynomial), then the dimension of its *eigenspace* (i.e. the subspace of vectors satisfying $\mathbf{A}v = \lambda v$) is also equal to m . This is not necessarily the case for a general matrix \mathbf{A} .
2. If $\lambda_i \neq \lambda_j$, this is simply a consequence of Proposition 2. Otherwise, find a basis of the eigenspace and turn it into an orthogonal basis using the Gram-Schmidt algorithm.
3. This is one is straightforward: we are simply saying that we can choose the vectors so that they have unit norm.

Spectral Decomposition iii

4. First, note that if Λ is a diagonal matrix with $\lambda_1, \dots, \lambda_n$ on the diagonal, and P is a matrix whose i -th column is v_i , then $\mathbf{A} = \sum_{i=1}^n \lambda_i v_i v_i^T$ is equivalent to

$$\mathbf{A} = P\Lambda P^T.$$

Then 4. is a consequence of the change of basis theorem: if we change the basis from the standard one to $\{v_1, \dots, v_n\}$, then \mathbf{A} acts by scalar multiplication in each direction, i.e. it is represented by a diagonal matrix Λ .



Examples i

We looked at

$$\mathbf{A} = \begin{pmatrix} 1 & 0.5 \\ 0.5 & 1 \end{pmatrix},$$

and determined that the eigenvalues were 1.5, 0.5, with corresponding eigenvectors $(1/\sqrt{2}, 1/\sqrt{2})$ and $(1/\sqrt{2}, -1/\sqrt{2})$.

Examples ii

```
v1 <- c(1/sqrt(2), 1/sqrt(2))  
v2 <- c(1/sqrt(2), -1/sqrt(2))
```

```
Lambda <- diag(c(1.5, 0.5))  
P <- cbind(v1, v2)
```

```
P %*% Lambda %*% t(P)
```

```
##      [,1] [,2]  
## [1,]  1.0  0.5  
## [2,]  0.5  1.0
```


Examples iii

Now let's look at a random matrix----

```
A <- matrix(rnorm(3 * 3), ncol = 3, nrow = 3)
```

Let's make it symmetric

```
A[lower.tri(A)] <- A[upper.tri(A)]
```

A

```
##           [,1]      [,2]      [,3]
## [1,] -0.59299188  0.01828677 -0.05900176
## [2,]  0.01828677  0.11304816 -0.70103784
## [3,] -0.05900176 -0.70103784  0.35028876
```

Examples iv

```
result <- eigen(A, symmetric = TRUE)
Lambda <- diag(result$values)
P <- result$vectors
```

```
P %*% Lambda %*% t(P)
```

```
##           [,1]      [,2]      [,3]
## [1,] -0.59299188  0.01828677 -0.05900176
## [2,]  0.01828677  0.11304816 -0.70103784
## [3,] -0.05900176 -0.70103784  0.35028876
```

Examples v

How to check if they are equal?

```
all.equal(A, P %*% Lambda %*% t(P))
```

```
## [1] TRUE
```

Positive-definite matrices

Let \mathbf{A} be a real symmetric matrix, and let $\lambda_1 \geq \dots \geq \lambda_n$ be its (real) eigenvalues.

1. If $\lambda_i > 0$ for all i , we say \mathbf{A} is *positive definite*.
2. If the inequality is not strict, if $\lambda_i \geq 0$, we say \mathbf{A} is *positive semidefinite*.
3. Similarly, if $\lambda_i < 0$ for all i , we say \mathbf{A} is *negative definite*.
4. If the inequality is not strict, if $\lambda_i \leq 0$, we say \mathbf{A} is *negative semidefinite*.

Note: If \mathbf{A} is *positive-definite*, then it is invertible!

Square Root Matrix i

- Let \mathbf{A} be a positive semidefinite symmetric matrix.
- By the Spectral Decomposition, we can write

$$\mathbf{A} = \mathbf{P}\mathbf{\Lambda}\mathbf{P}^T.$$

- Since \mathbf{A} is positive-definite, we know that the elements on the diagonal of $\mathbf{\Lambda}$ are positive.
- Let $\mathbf{\Lambda}^{1/2}$ be the diagonal matrix whose entries are the square root of the entries on the diagonal of $\mathbf{\Lambda}$.
- For example:

$$\mathbf{\Lambda} = \begin{pmatrix} 1.5 & 0 \\ 0 & 0.5 \end{pmatrix} \Rightarrow \mathbf{\Lambda}^{1/2} = \begin{pmatrix} 1.2247 & 0 \\ 0 & 0.7071 \end{pmatrix}.$$

Square Root Matrix ii

- We define the square root $\mathbf{A}^{1/2}$ of \mathbf{A} as follows:

$$\mathbf{A}^{1/2} := P\Lambda^{1/2}P^T.$$

- *Check:*

$$\begin{aligned}\mathbf{A}^{1/2}\mathbf{A}^{1/2} &= (P\Lambda^{1/2}P^T)(P\Lambda^{1/2}P^T) \\ &= P\Lambda^{1/2}(P^TP)\Lambda^{1/2}P^T \\ &= P\Lambda^{1/2}\Lambda^{1/2}P^T \quad (P \text{ is orthogonal}) \\ &= P\Lambda P^T \\ &= \mathbf{A}.\end{aligned}$$

Power method

Introduction to numerical algebra

- As presented in these notes, we can find the eigenvalue decomposition by
 1. Finding the roots of a degree n polynomial.
 2. For each root, find the solutions to a system of linear equations.
- Problem: no exact formula for roots of a generic polynomial when $n > 4$.
 - So we need to find approximate solutions
- Other problem: approximation errors for eigenvalues propagate to eigenvectors
- **Need more stable algorithms**
- This is what numerical algebra is about. For a good reference, I recommend *Matrix Computations* by Golub and Van Loan.

Power Method i

- We'll discuss one approach to finding the leading eigenvector, i.e. the eigenvector corresponding to the largest eigenvalue (in absolute value).
- **Note:** We have to assume that the largest eigenvalue (in absolute value) is unique.
- *Algorithm:*
 1. Let v_0 be an initial vector with unit norm.
 2. At step k , define

$$v_{k+1} = \frac{\mathbf{A}v_k}{\|\mathbf{A}v_k\|},$$

where $\|v\|$ is the norm of the vector v .

Power Method ii

3. Then the sequence v_k converges to the desired eigenvector.
4. The corresponding eigenvalue is defined by

$$\lambda = \lim_{k \rightarrow \infty} \frac{v_k^T \mathbf{A} v_k}{v_k^T v_k}.$$

- Comment: unless v_0 is orthogonal to the eigenvector we are looking for, we have theoretical guarantees of convergence.
 - In practice, we can pick v_0 randomly, since the probability a random vector is orthogonal to the eigenvector is zero.

Example i

```
set.seed(123)

A <- matrix(rnorm(3*3), ncol = 3)
# Make A symmetric
A[lower.tri(A)] <- A[upper.tri(A)]

# Set initial value
v_current <- rnorm(3)
v_current <- v_current/norm(v_current, type = "2")
```

Example ii

```
# We'll perform 100 iterations
for (i in seq_len(100)) {
  # Save result from previous iteration
  v_previous <- v_current
  # Compute matrix product
  numerator <- A %*% v_current
  # Normalize
  v_current <- numerator/norm(numerator, type = "2")
}

v_current
```

Example iii

```
##           [,1]
## [1,] -0.3318109
## [2,]  0.5345952
## [3,]  0.7772448
```

Corresponding eigenvalue

```
num <- t(v_current) %*% A %*% v_current
denom <- t(v_current) %*% v_current
num/denom
```

```
##           [,1]
## [1,] -1.75374
```

Example iv

```
# CHECK results
```

```
result <- eigen(A, symmetric = TRUE)
```

```
result$values[which.max(abs(result$values))]
```


```
## [1] -1.75374
```

```
result$vectors[,which.max(abs(result$values))]
```

```
## [1] 0.3318109 -0.5345952 -0.7772448
```

- Note that we did not get the same eigenvector: they differ by -1.

Visualization



Blue is the objective; the sequence goes from green to red.